Whirlpool

Data Acquisition using N-node Distributed Web Crawler

Rihan Pereira, MSCS

Advisor: Dr. Michael Soltys
Department of Computer Science
MSCS Graduate 2018-2019

November 27, 2019

1. Motivation & Contribution
2. Crawler characteristics & history
3. Mercator 1999 (Heydon & Najork)
4. Software Design Principles
5. Whirlpool: Event-driven architecture
6. Whirlpool: Parser
7. Whirlpool: Near-Deduplication
8. Whirlpool: Distributed Crawling
9. Whirlpool: Operations
10. Future work

Motivation & Contribution

Motivation

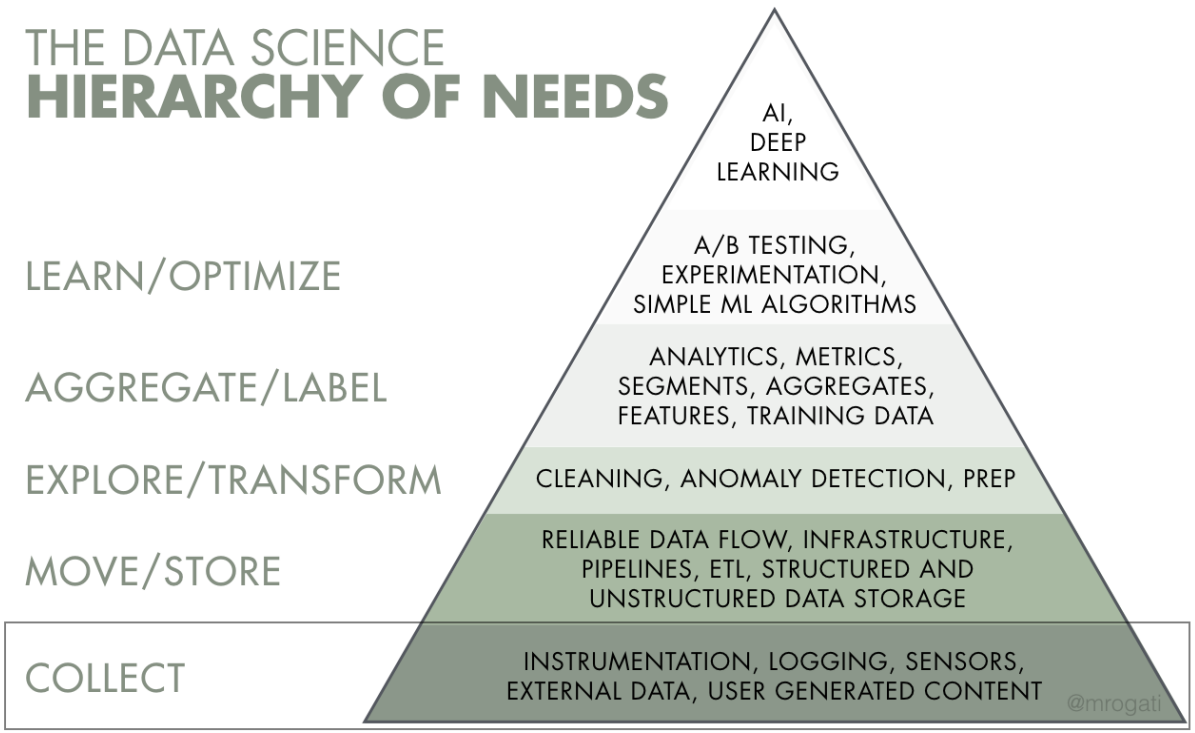
THE DATA SCIENCE
HIERARCHY OF NEEDS
AI, DEEP LEARNING
LEARN/OPTIMIZE
A/B TESTING, EXPERIMENTATION, SIMPLE ML ALGORITHMS
AGGREGATE/LABEL
ANALYTICS, METRICS, SEGMENTS, AGGREGATES, FEATURES, TRAINING DATA
EXPLORE/TRANSFORM
CLEANING, ANOMALY DETECTION, PREP
MOVE/STORE
RELIABLE DATA FLOW, INFRASTRUCTURE, PIPELINES, ETL, STRUCTURED AND UNSTRUCTURED DATA STORAGE
COLLECT
INSTRUMENTATION, LOGGING, SENSORS, EXTERNAL DATA, USER GENERATED CONTENT
@mrogati

Motivation

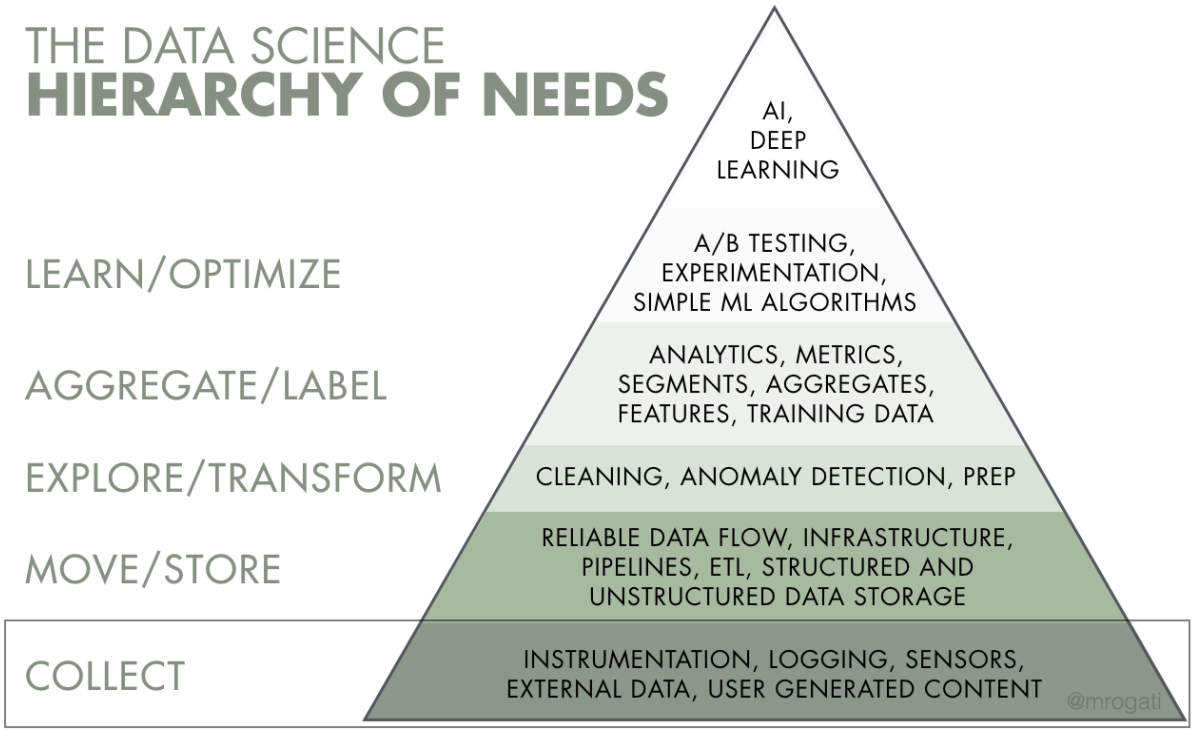

Self-actualization (AI) is great, but you first need food, water, and shelter (data literacy, collection, and infrastructure).”

Contributions

to be completed

Crawler characteristics & history

Coverage & Freshness

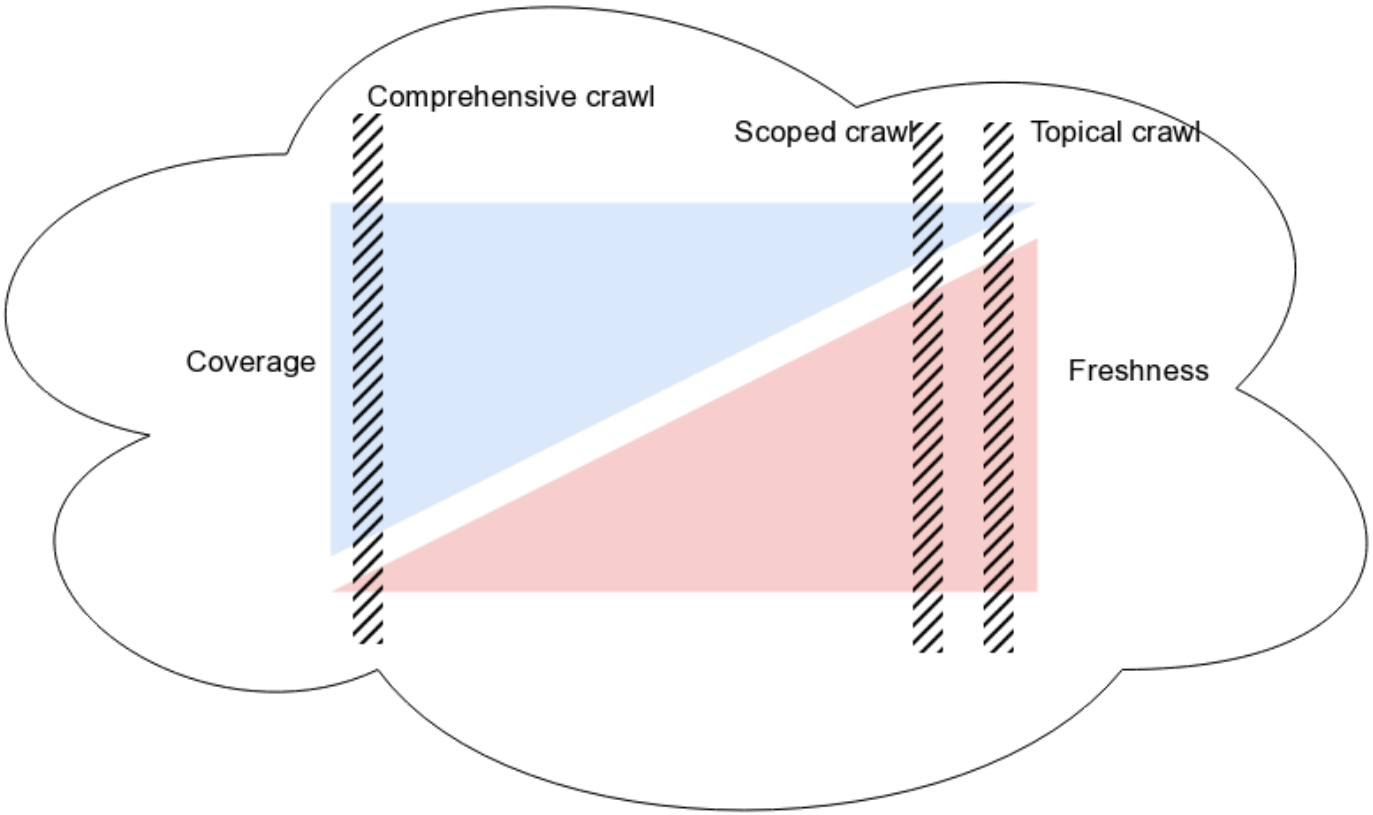

Web crawlers (1990 - 2019)

to add something

Mercator 1999 (Heydon & Najork)

basic crawling algorithm

to add content

Mercator background

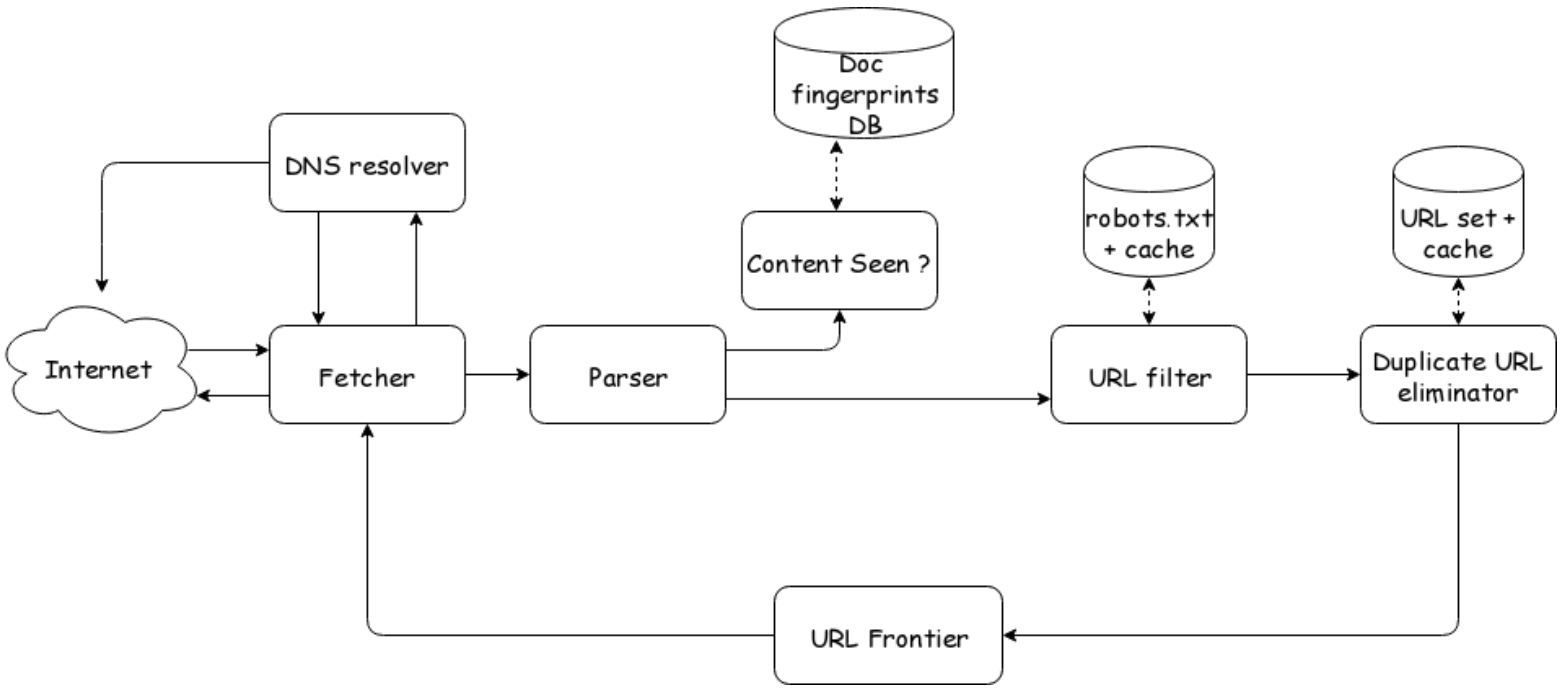

Figure: Mercator building blocks (Heydon & Najork)

URL Frontier Scheme

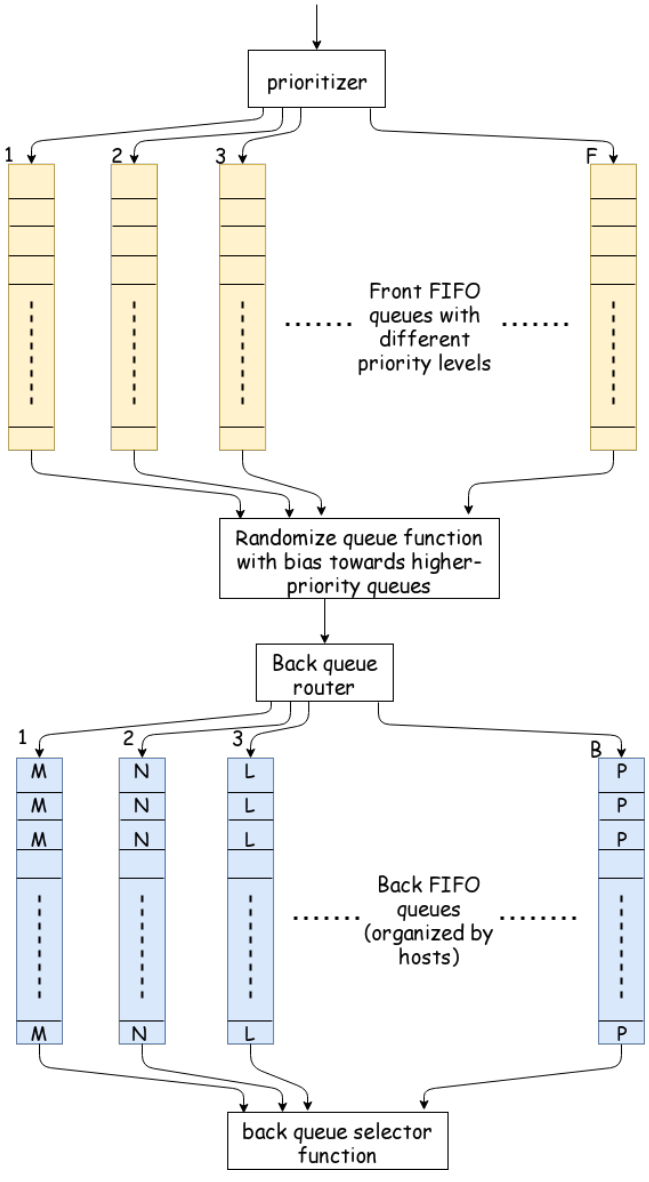

Front queue (Frontier Queue)

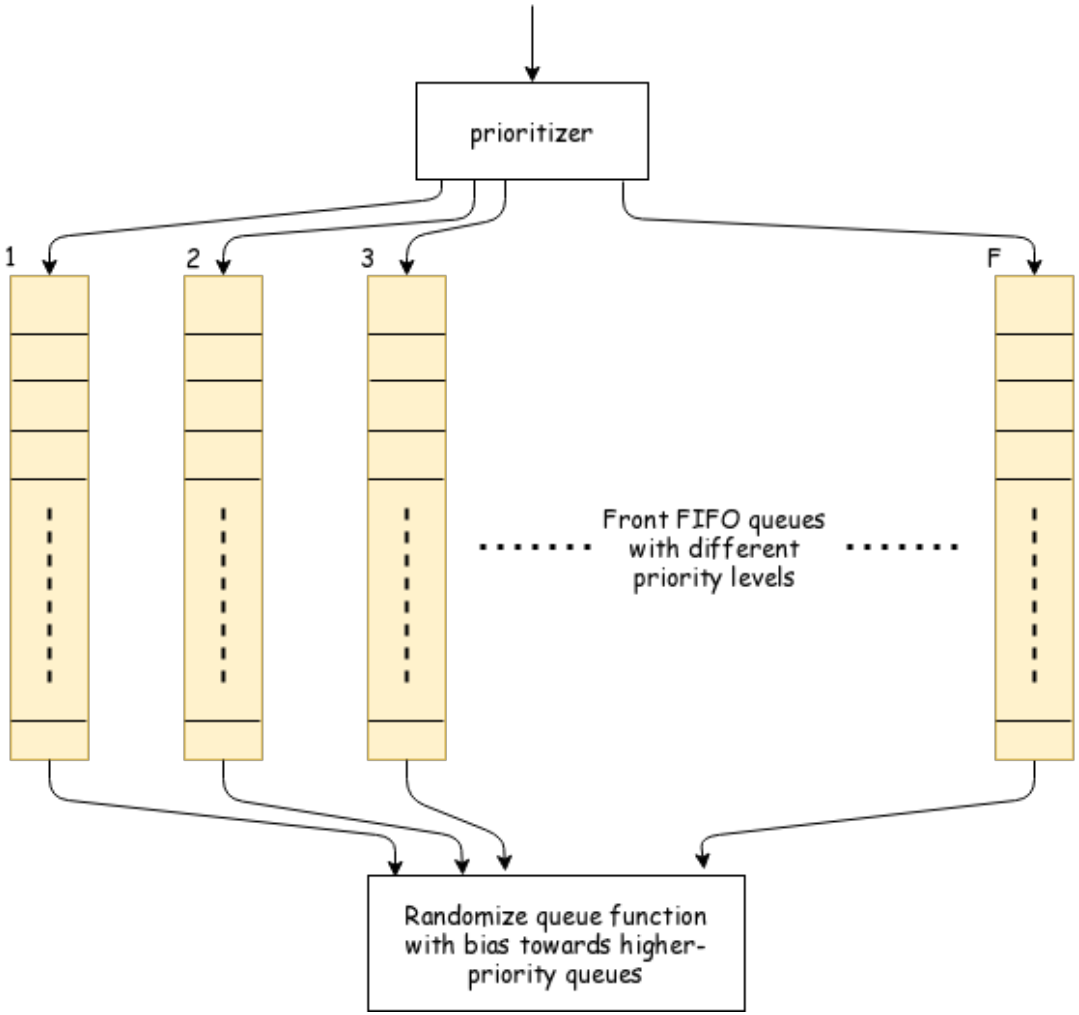

Back queue (Frontier Queue)

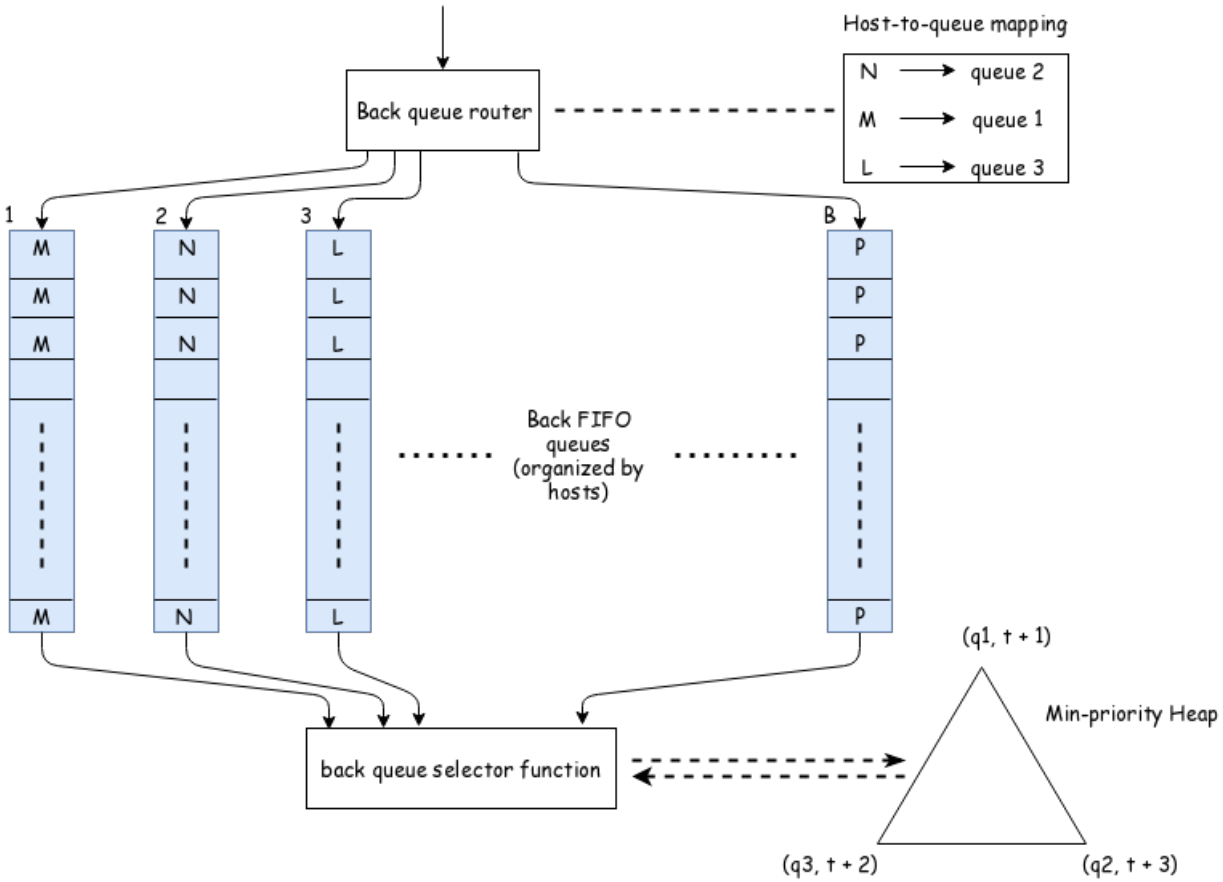

Software Design Principles

Designing scalable systems

Designing scalable systems

- Adding identical copies of components

Designing scalable systems

- Adding identical copies of components
- Functional partitioning

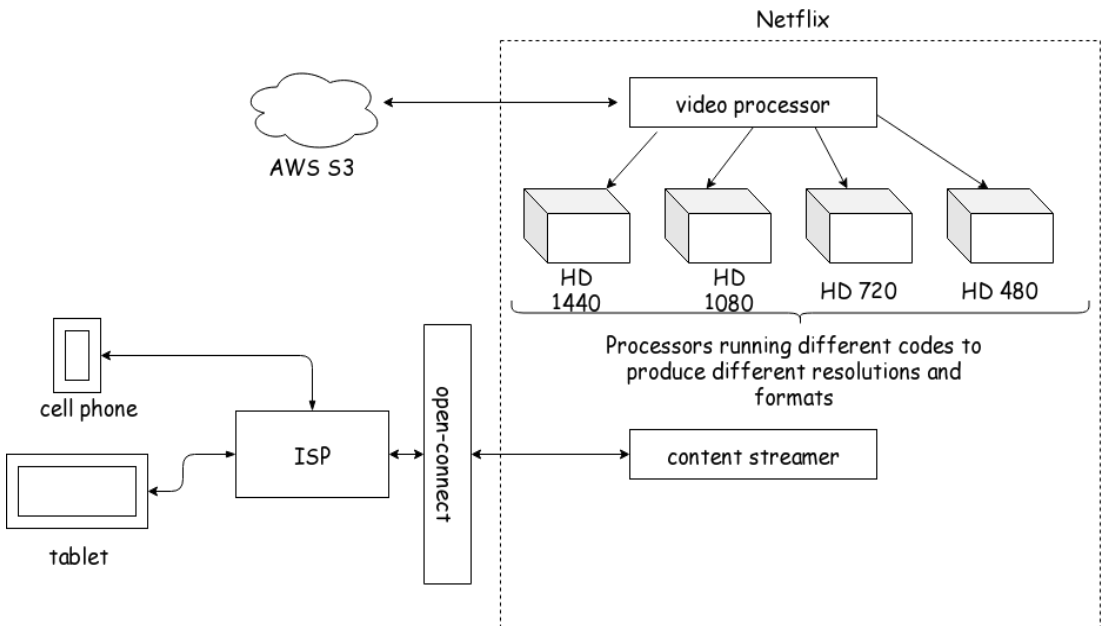

Designing scalable systems

- Adding identical copies of components
- Functional partitioning

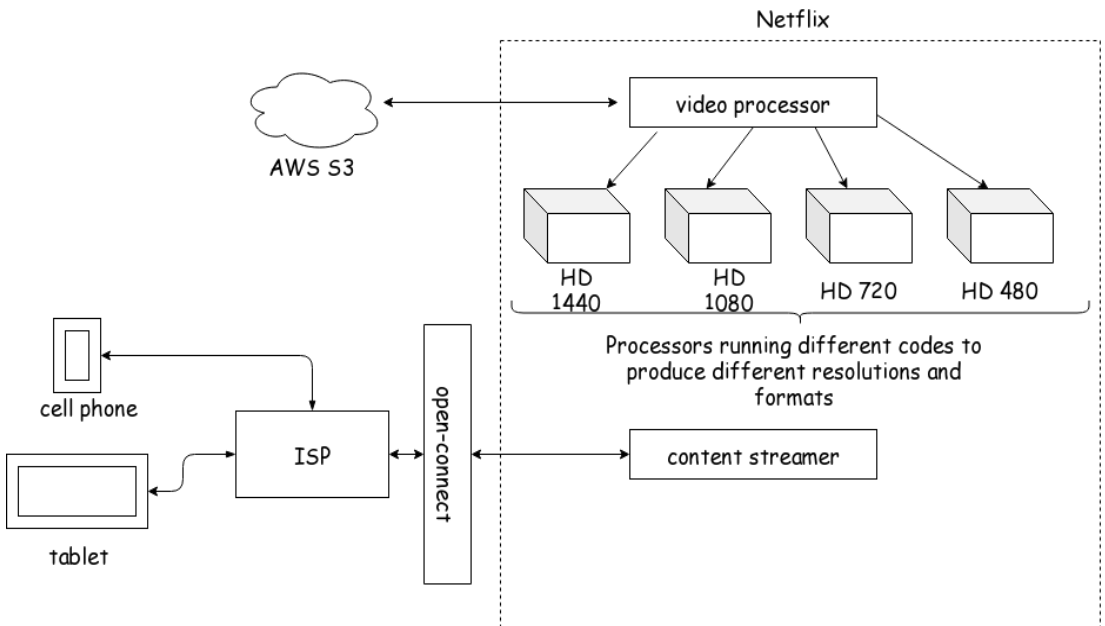

- Data partitioning

State Management

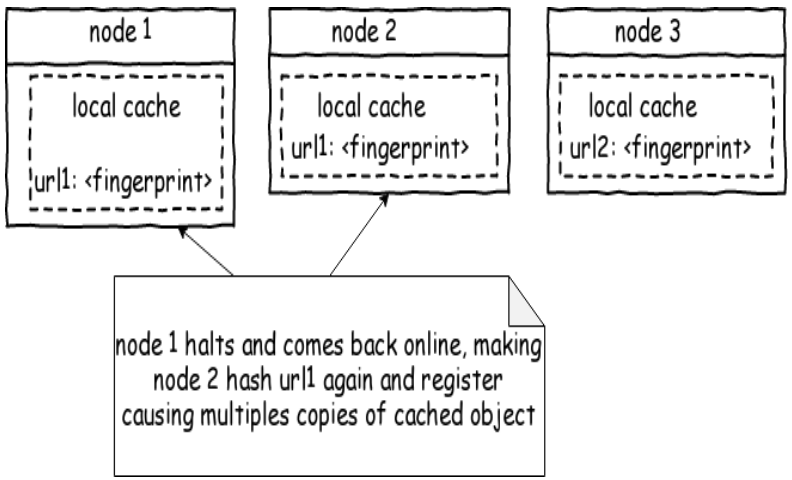

Figure: identical copies of same cached object

State Management

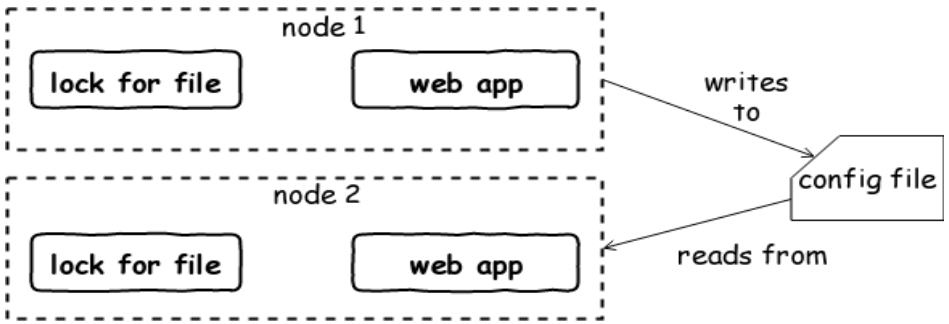

Figure: Using local locks to access shared resources

State Management

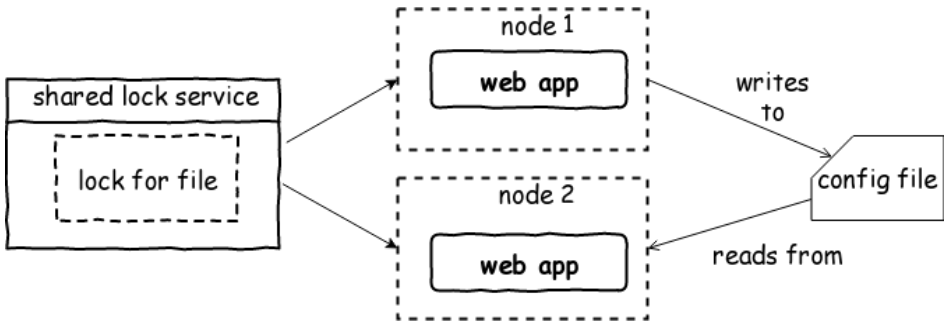

Figure: using shared locks to access shared resources

Whirlpool: Event-driven architecture

Message Queue(MQ)

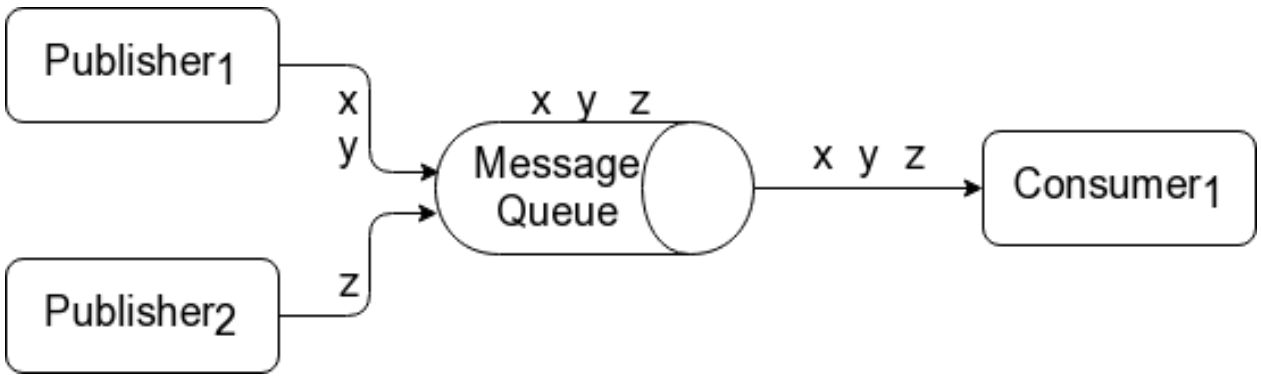

Message Queue(MQ)

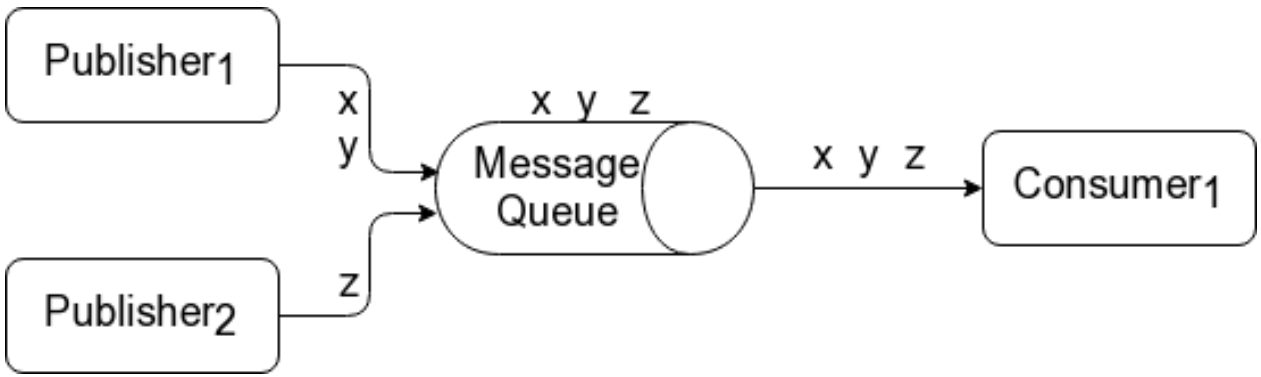

- decoupling

Message Queue(MQ)

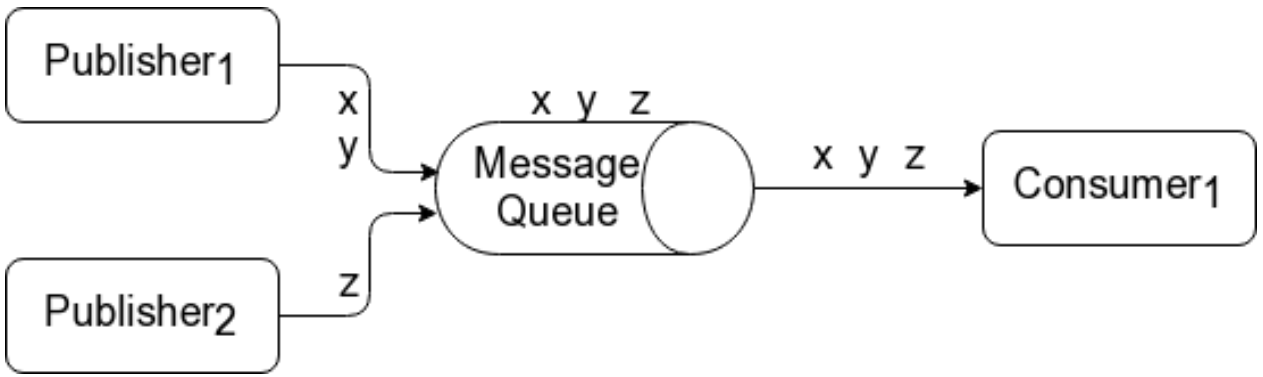

- decoupling
- scaled independently

Message Queue(MQ)

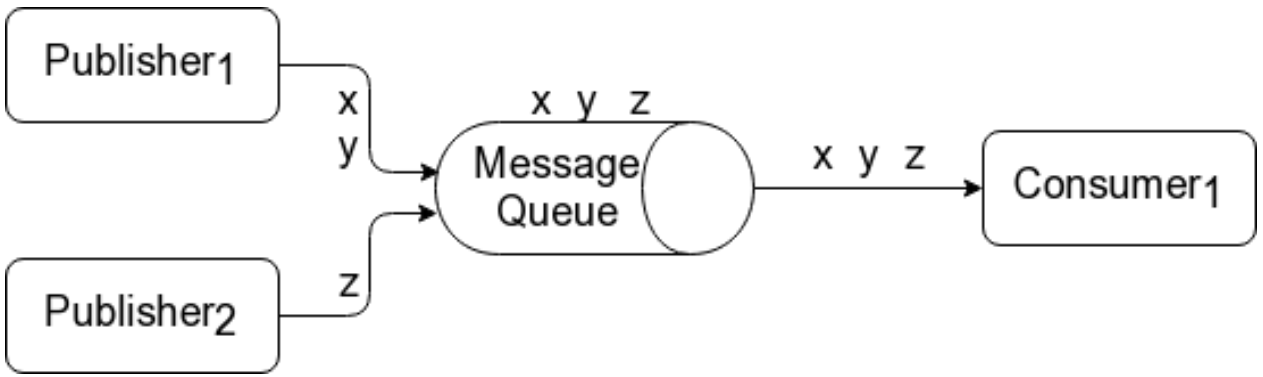

- decoupling
- scaled independently
- balancing traffic

Message Queue(MQ)

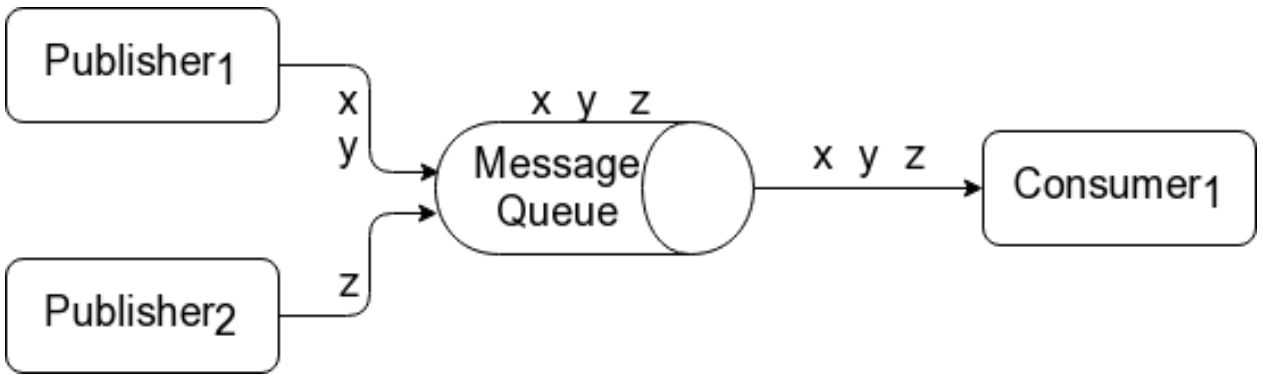

- decoupling
- scaled independently
- balancing traffic
- fault-tolerance

MQ: Routing mechanisms

MQ: Routing mechanisms

- Direct Worker Queue Data Flow

MQ: Routing mechanisms

- Direct Worker Queue Data Flow
- Fanout

MQ: Routing mechanisms

- Direct Worker Queue Data Flow
- Fanout
- Topic

MQ: Routing mechanisms

- Direct Worker Queue Data Flow
- Fanout
- Topic
- Header

Direct Worker Queue Data Flow

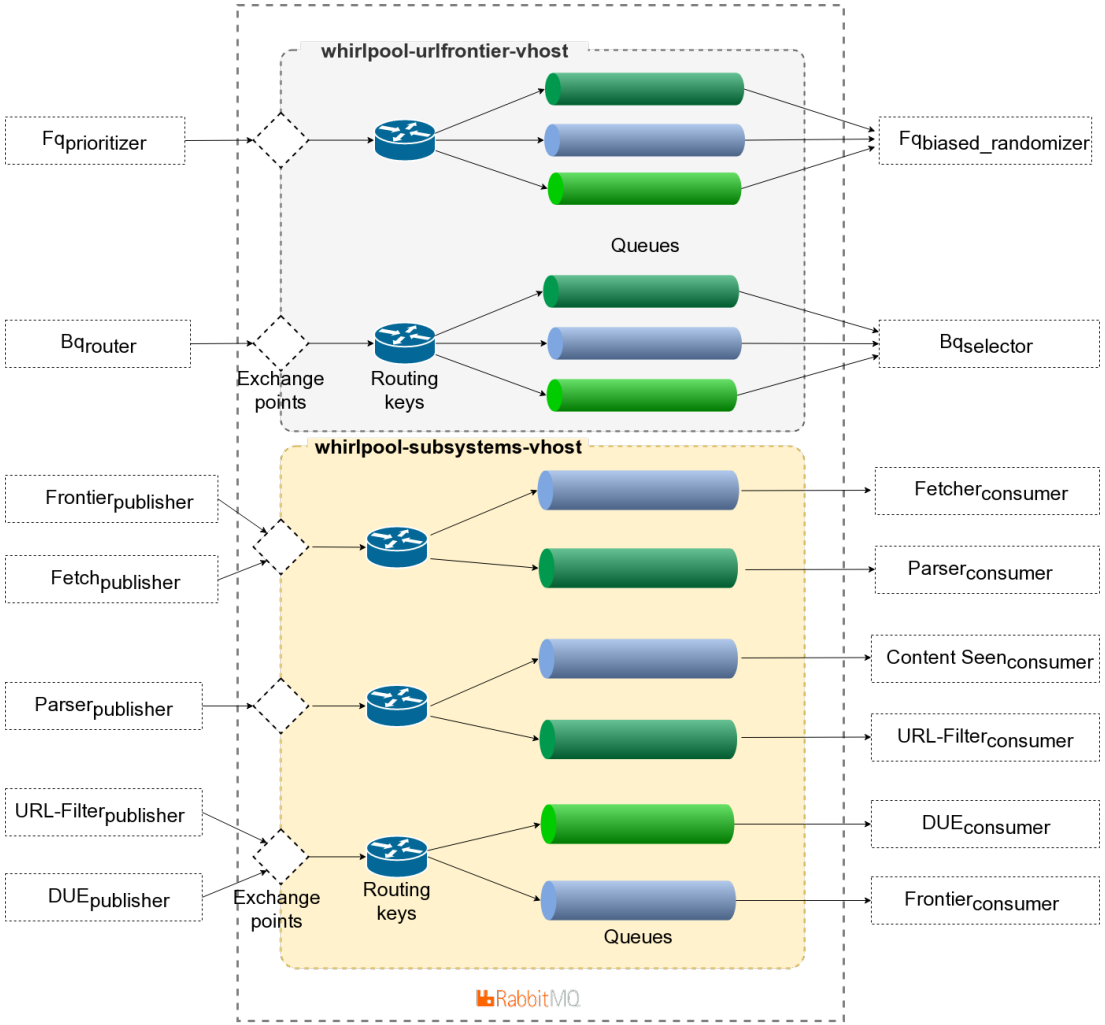

RabbitMQ: Message bus

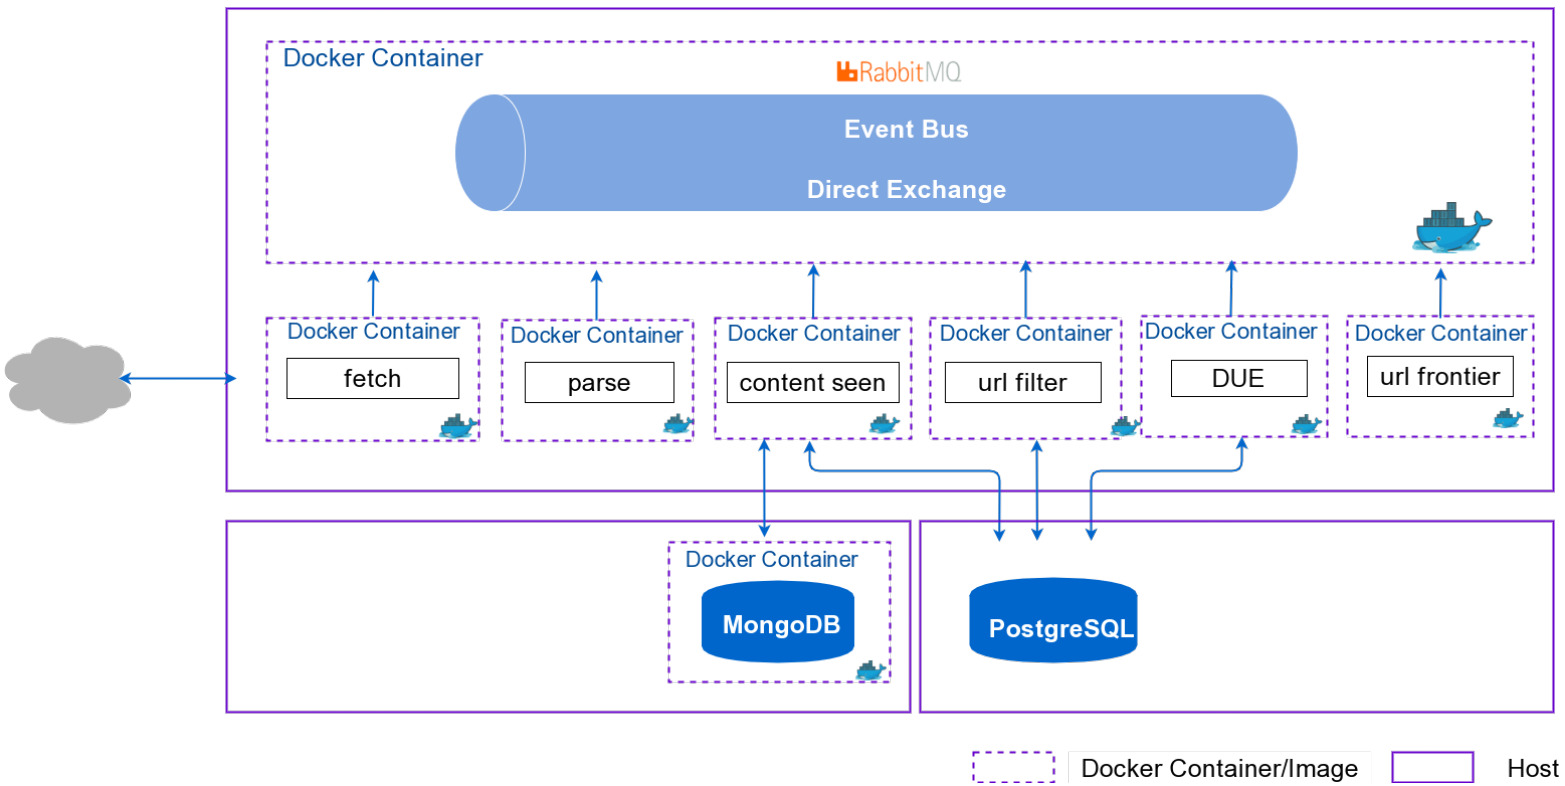

development vs. production docker containers

things to add

Whirlpool: Parser

Parser

to add something

Whirlpool: Near-Deduplication

Dedupe

to add something

Whirlpool: Distributed Crawling

Dist. crawl

to add something

Whirlpool: Operations

From 10,000 ft.

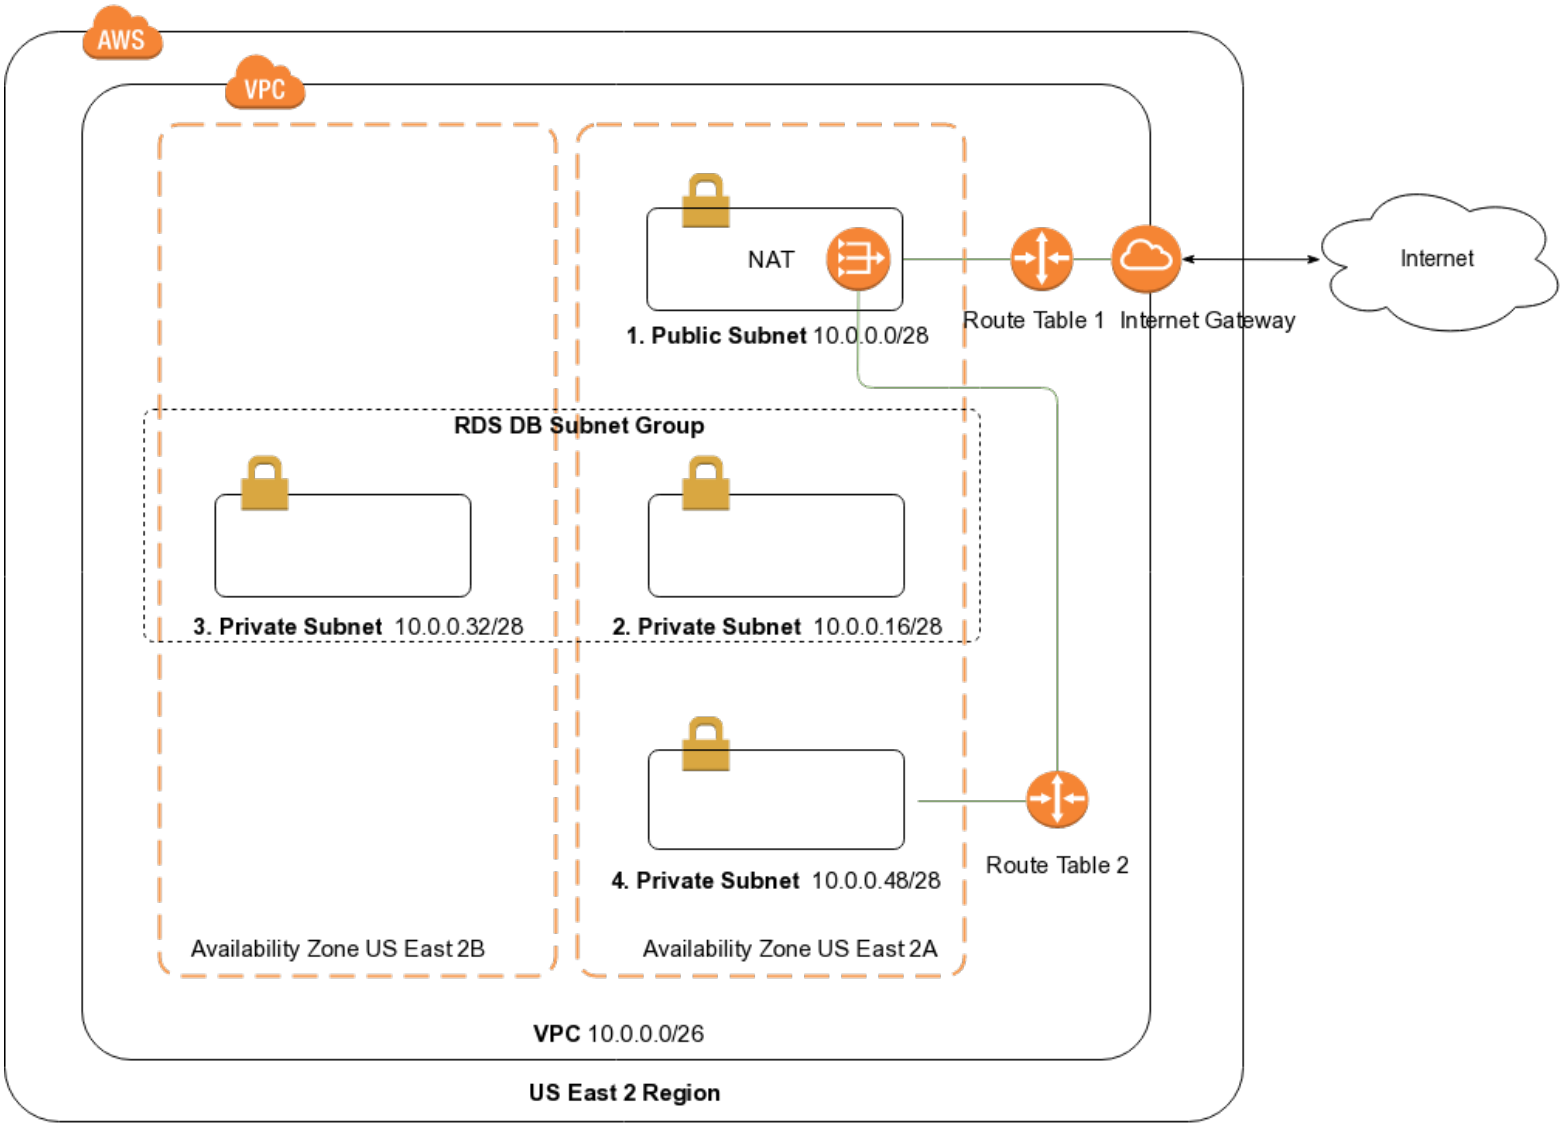

From 5,000 ft.

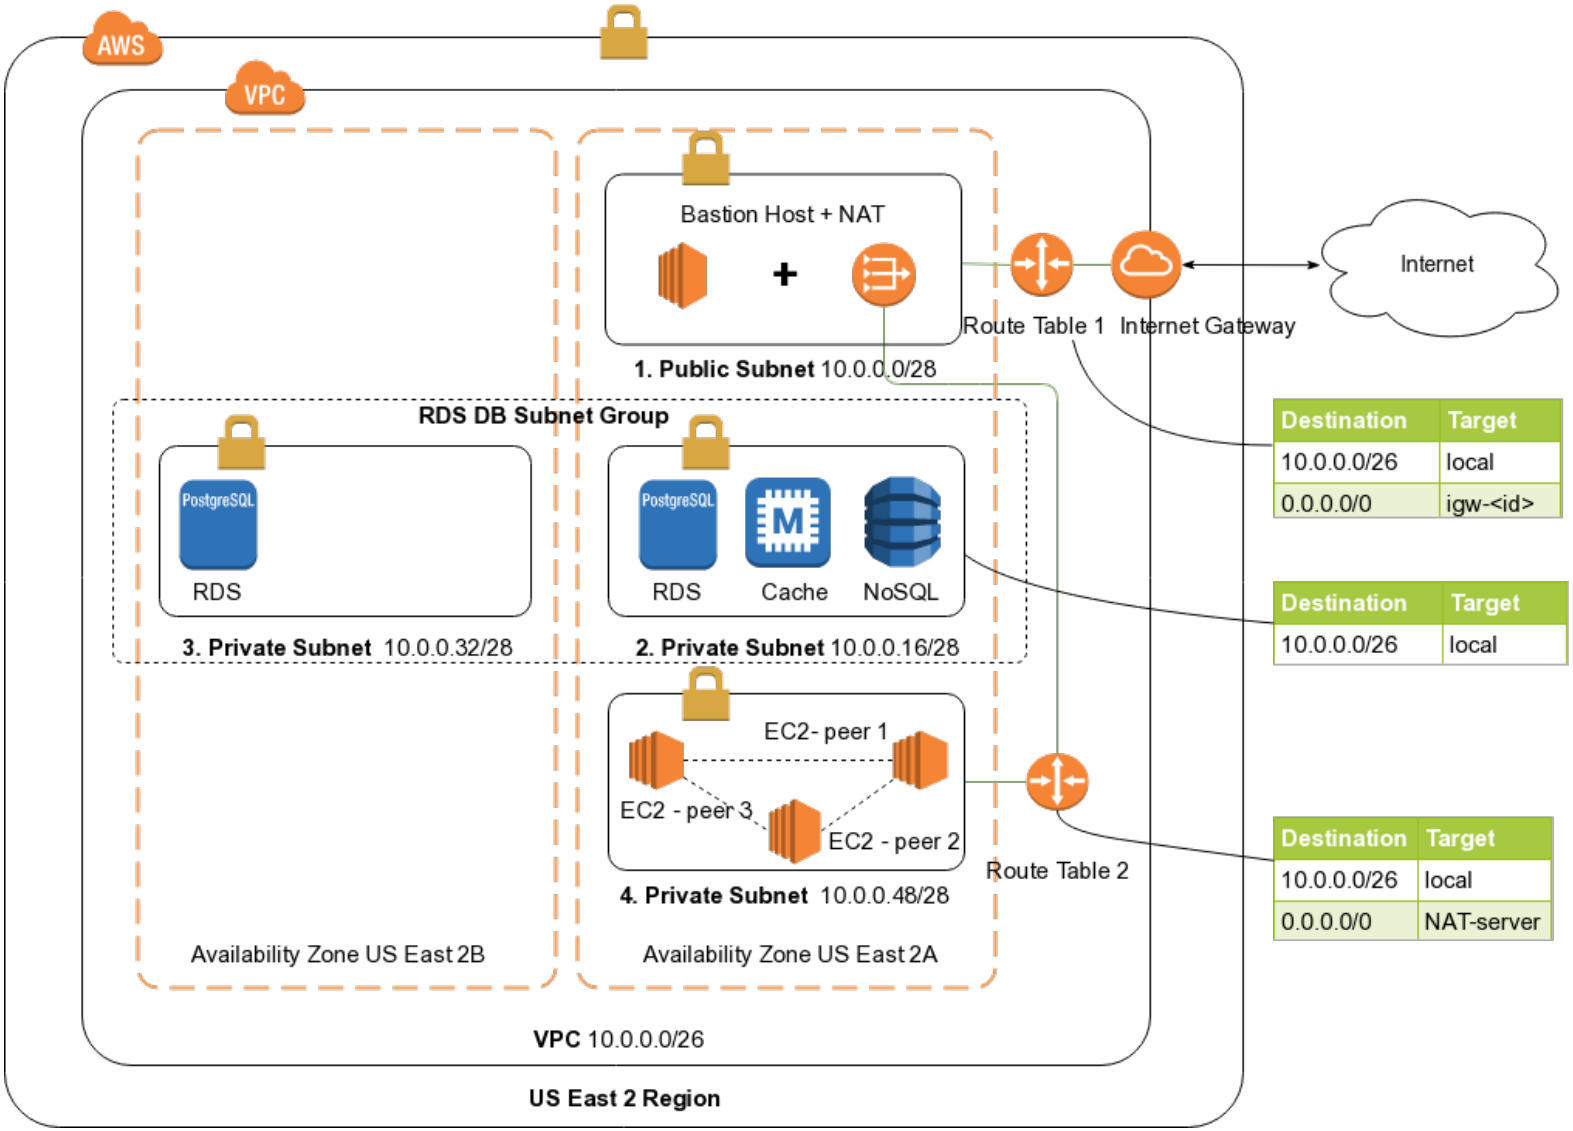
AWS
VPC
Bastion Host + NAT
1. Public Subnet 10.0.0.0/28
Route Table 1
Internet Gateway
Internet
RDS DB Subnet Group
PostgreSQL
RDS
3. Private Subnet 10.0.0.32/28
PostgreSQL
RDS
Cache
NoSQL
2. Private Subnet 10.0.0.16/28
EC2- peer 1
EC2 - peer 3
EC2 - peer 2
4. Private Subnet 10.0.0.48/28
Route Table 2
Availability Zone US East 2B
Availability Zone US East 2A
VPC 10.0.0.0/26
US East 2 Region
Destination Target
10.0.0.0/26 local
0.0.0.0/0 igw-<id>
Destination Target
10.0.0.0/26 local
Destination Target
10.0.0.0/26 local
0.0.0.0/0 NAT-server

Future work

future to do

to add something

Thank you! Questions ?